

# Reinforcement Learning

## Gymnasium

Jordi Casas Roma

[jordi.casas.roma@uab.cat](mailto:jordi.casas.roma@uab.cat)

September 12, 2025



# Table of Contents

## Introduction

### Basic structure of an RL scenario

- Description

- Basic operating diagram

## The Gymnasium API

- Space

- Environment

- Wrappers

## Classic RL Examples

- CartPole

- FrozenLake

## Bibliography


# Introduction

## Gymnasium

- ▶ An **API standard** for reinforcement learning with a diverse collection of reference environments.
- ▶ <https://gymnasium.farama.org/index.html>
- ▶ Version 1.2.0



An API standard for reinforcement learning with a diverse collection of reference environments



Gymnasium

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Introduction

## Description


In a reinforcement learning environment there are **two main entities**:

- ▶ The **agent** is an entity that **plays an active role** in the environment, making decisions about the actions it will execute in the environment.
  - ▶ From the most pragmatic point of view, the agent is an application (code fragment) that implements a certain policy to execute certain actions in the environment.
  - ▶ Basically, this policy decides what action is required at any given time based on observations provided by the environment.
- ▶ The **environment** is a model external to the agent that has the **responsibility of providing observations and rewards to the agent's actions**.
  - ▶ The environment changes its state based on the agent's actions.

# Introduction

## Description

- We can see that for each **action** ( $A_t$ ) that the agent performs at time instant  $t$ , the environment updates its state and returns:
  1. the **reward** associated with the action ( $R_{t+1}$ ), and
  2. the **observation** associated with the new state ( $O_{t+1}$ )



# Table of Contents

Introduction

**Basic structure of an RL scenario**

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Introduction

## Basic operating diagram

### Main code block:

- ▶ The classes corresponding to the **environment (Environment)** and the **agent (Agent)** are initialized.
- ▶ Once the classes are initialized, the agent iterates indefinitely using the **step() method** until the episode ends.

```
1 if __name__ == "__main__":
2     env = Environment()
3     agent = Agent()
4
5     while not env.is_done():
6         agent.step(env)
7
8     print("Total reward got: {}".format(agent.total_reward))
```

# Introduction

## The Agent

The **agent** class:

1. The `__init__()` method: initializes the internal variables of the agent.
2. The `step()` method: execute the actions at time  $t$ .

```
1 class Agent:  
2     def __init__(self):  
3         # Agent init  
4         self.total_reward = 0.0  
5  
6     def step(self, env: Environment):  
7         current_obs = env.get_observation()  
8         actions = env.get_actions()  
9         action = <SELECT desired action>  
10        reward = env.action(action)  
11        self.total_reward += reward
```

# Introduction

## The Agent

The **agent** class:

1. The `__init__()` method: initializes the internal variables of the agent.
2. The `step()` method: execute the actions at time  $t$ .

```
1 class Agent:  
2     def __init__(self):  
3         # Agent init  
4         self.total_reward = 0.0  
5  
6     def step(self, env: Environment):  
7         current_obs = env.get_observation()  
8         actions = env.get_actions()  
9         action = <SELECT desired action>  
10        reward = env.action(action)  
11        self.total_reward += reward
```

- What does the `step()` method do?

# Introduction

## Basic operating diagram

What does the `step()` method do?

1. Obtains the **value of the observation** of the current state, which is provided by the environment.
2. Obtains the **values of possible actions**, which is also provided by the environment.
3. Next, with the previous information, the agent **must select the action** that it wants to perform.
  - ▶ At this point is where the agent implements the selected **method to choose the action** to perform among the set of possible actions.
  - ▶ This **set can be limited or unlimited**, depending on the characteristics of the environment
4. Finally, it **communicates the action to the environment**, which returns the associated **reward**.

# Introduction

## The Environment

The **environment** class:

1. The environment must provide several methods to report **current observations**, **possible actions**, the **status of the episode**.
2. The **action()** method, which applies a certain action to the environment and updates it, providing the indicated **reward**.

```
1 class Environment:  
2     def __init__(self):  
3         # Environment initialization  
4  
5     def get_observation(self):  
6         # Returns the observation values of the current state  
7  
8     def get_actions(self):  
9         # Returns the list of possible actions  
10  
11    def is_done(self):  
12        # Returns a boolean indicating whether the episode has ended  
13  
14    def action(self, action):  
15        # Apply the indicated action in the environment  
16        # Returns the associated reward
```

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# The Gymnasium API

## Introduction

### Introduction:

1. **Gymnasium** is a library for **developing and comparing reinforcement learning algorithms**.
2. Gymnasium makes **no assumptions about the structure of the agent you want to implement**, and is compatible with any numerical computing library, such as TensorFlow, Theano, or PyTorch.
3. Gymnasium offers a **collection of test environments** that can be used to **train, evaluate and compare** reinforcement learning algorithms.
4. These environments have a **shared interface** (API) that allows general algorithms to be developed.

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake


Bibliography

# The Gymnasium API

## Space

### Space:

1. The Space class allows you to define what the `actions` and `states` will be like, defining the type and range of allowed values.
2. There are `multiple subclasses` that implement the abstract class Space.
3. Any environment contains two members of type Space, called `action_space` and `observation_space`.



# The Gymnasium API

## Space

The methods of **Space** class:

- ▶ The superclass Space includes two relevant methods:
  - ▶ `sample()`, which returns a random sample of the space.
  - ▶ `contains(x)`, which checks whether the argument `x` belongs to the domain space.
- ▶ Both methods are abstract and are re-implemented in each of their subclasses.

# The Gymnasium API

## Space

The methods of **Discrete** subclass:

- ▶ The class **Discrete** represents a set of mutually exclusive elements, numbered from 0 to  $n - 1$ .
- ▶ Its main attribute ( $n$ ) is a count of the elements it describes.

# The Gymnasium API

## Space

The methods of **Discrete** subclass:

- ▶ The class **Discrete** represents a set of mutually exclusive elements, numbered from 0 to  $n - 1$ .
- ▶ Its main attribute ( $n$ ) is a count of the elements it describes.

### Movement on a chess board

For example, `Discrete(n=4)` can be used to define an **action space** that represents movement in four directions, where the value 0 indicates move left, 1 moves right, 2 moves up, and 3 indicates moves down.

# The Gymnasium API Space

The methods of **Box** subclass:

- ▶ The class Box represents an *n-dimensional tensor* of rational numbers with the interval [low, high].
- ▶ This class is *widely used*, since it allows representing a large amount and variety of information of a very different nature.

# The Gymnasium API

## Space

The methods of **Box** subclass:

- ▶ The class Box represents an *n-dimensional tensor* of rational numbers with the interval [low, high].
- ▶ This class is *widely used*, since it allows representing a large amount and variety of information of a very different nature.

### Acceleration pedal

For example, an accelerator pedal can be represented by a *single value between 0.0 and 1.0*, indicating the active percentage of the accelerator.

```
Box(low=0.0, high=1.0, shape=(1,), dtype=np.float32)
```

- ▶ The *shape* argument indicates that this is a *tuple of length 1 with a single value*.
- ▶ The *range* of this value is determined by the *low* and *high* values.
- ▶ The *dtype* parameter specifies the *type* of the variable, which in this particular case is a 32-bit NumPy float.

# The Gymnasium API

## Space

The methods of **Tuple** subclass:

- ▶ The Tuple class allows us to **combine several instances** of the Space class to create action and observation spaces of variable complexity and a high degree of adaptation to any environment.

# The Gymnasium API

## Space

The methods of **Tuple** subclass:

- ▶ The Tuple class allows us to **combine several instances** of the Space class to create action and observation spaces of variable complexity and a high degree of adaptation to any environment.

### Driving a vehicle

The vehicle has several controls, including the **angle of the steering wheel**, **the position of the brake pedal**, and the **position of the accelerator pedal**. These three controls can be specified by **three float values** in a single Box instance.

However, the vehicle requires additional discrete controls: **turn signal** (which could be off, left or right) and the **horn** (on or off).

To combine all of this into an action space specification class:

```
Tuple(spaces=(Box(low=-1.0, high=1.0, shape=(3,),  
dtype=np.float32), Discrete(n=3), Discrete(n=2)))
```

# The Gymnasium API

## Space

### Atari Games example

Suppose we want to represent an observation of an **Atari game screen**, which is an **image** of size  $210 \times 160$  in RGB (red, green and blue) using **one byte** to encode each colour, that is, 256 possible values.

**How can we encode this information?**

# The Gymnasium API

## Space

### Atari Games example

Suppose we want to represent an observation of an **Atari game screen**, which is an **image** of size  $210 \times 160$  in RGB (red, green and blue) using **one byte** to encode each colour, that is, 256 possible values.

#### How can we encode this information?

##### Atari game

We can represent this information using the following encoding:

```
Box(low=0, high=255, shape=(210, 160, 3), dtype=np.uint8)
```

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

**Environment**

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# The Gymnasium API

## Environment

**Environment:** The main goal of Gymnasium is to provide a large collection of environments with a common interface to allow different implementations to be evaluated and compared.

The environment is represented by the class `Env`, which has the following **methods**:

- ▶ `action_space`: Provides a specification for the actions allowed in the environment.
- ▶ `observation_space`: Specifies the observations provided by the environment.
- ▶ `reset()`: This method resets the environment to its initial state, returning the initial observation vector.
- ▶ `step()`: This method allows us to execute a certain action and returns information about its result, that is, the next observation, the reward and the end of the episode indicator.
- ▶ `render()`: This method allows us to obtain the observation in a human-friendly form.

# The Gymnasium API

## Environment

The **Env** class:

- ▶ The `step()` method is the central method in the environment's functionality. When this method is called, **several actions are performed**:
  1. Tell the environment what `action` we will execute in the next step.
  2. Get the new `observation` of the environment after this action.
  3. Get the `reward` generated by the action.
  4. Get a boolean indicating whether the `episode has ended`.
- ▶ The method has a **single main parameter**, which is the action that the agent wants to perform.

# The Gymnasium API

## Environment

The method `step()` of `Env` class:

The method returns a **five-element tuple** with the following types and meanings:

1. `observation`: NumPy vector or a matrix with the data of the next observation.
2. `reward`: Reward value, usually as a decimal (float) value.
3. `terminated`: Boolean flag, whether the agent reaches the terminal state.
4. `truncated`: Boolean flag, whether the truncation condition is satisfied (time limit, etc.).
5. `info`: additional information about the environment. In general, this value is usually ignored, although it can provide interesting information in certain cases.

# The Gymnasium API

## Environment

To instantiate the Env class, you must use the `make()` function of the Gym API.

1. This function has a **single argument** that must be indicated: the **name of the environment** to be created, which takes the form **EnvironmentName-vN**, where N is the number used to distinguish between different versions of the same environment (when, for example, some bugs are fixed or other important changes are made).

# The Gymnasium API

## Environment

To instantiate the Env class, you must use the `make()` function of the Gym API.

1. This function has a **single argument** that must be indicated: the **name of the environment** to be created, which takes the form **EnvironmentName-vN**, where N is the number used to distinguish between different versions of the same environment (when, for example, some bugs are fixed or other important changes are made).

For example, to create the environment “CartPole-v0” the following code must be executed:

```
1 import gym  
2  
3 env = gym.make('CartPole-v0')
```

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# The Gymnasium API

## Wrappers

### Wrappers:

- ▶ Wrappers are a convenient way **to modify an existing environment** without having to alter the underlying code directly.
- ▶ Wrappers can also be **chained** to combine their effects.

### Wrapper example

For example, when using image-based inputs, it is common to use a range of values [0,1] instead of [0,255], which are the standard values in RGB images.

# The Gymnasium API

## Wrappers


The **wrappers'** hierarchy:

- ▶ ObservationWrapper is used to modify the **observations** returned by the environment.
- ▶ RewardWrapper is used to modify the **rewards** returned by the environment.
- ▶ ActionWrapper allows you to modify the **actions** sent to the environment.

# The Gymnasium API

## Wrappers

The **wrappers'** hierarchy:



# The Gymnasium API

## Wrappers

In order to wrap an environment, we must:

1. First, **initialize** a base environment.
2. Then, we can pass this environment along with optional parameters to the **wrapper's constructor**.

```
1 import gymnasium as gym
2 from gymnasium.wrappers import RescaleAction
3
4 base_env = gym.make("Hopper-v4")
5 base_env.action_space
6
7 >> Box(-1.0, 1.0, (3,), float32)
8
9 wrapped_env = RescaleAction(base_env, min_action=0, max_action=1)
10 wrapped_env.action_space
11
12 >> Box(0.0, 1.0, (3,), float32)
```

# The Gymnasium API

## Wrappers

The implementation of the Wrapper class requires defining a `__init__` method that accepts the environment as a parameter.

```
1 class BasicWrapper(gym.Wrapper):
2     def __init__(self, env):
3         super().__init__(env)
4         self.env = env
5
6     def step(self, action):
7         next_state, reward, terminated, truncated, info = self.env.
8         step(action)
9         # modifications
10        return next_state, reward, terminated, truncated, info
```

And the new class `BasicWrapper` is instantiated as follows:

```
1 env = BasicWrapper(gym.make("EnvironmentName-vN"))
```

# The Gymnasium API

## Wrappers

Some **interesting** wrappers (among many others):

- ▶ RecordVideo: This wrapper records **videos** of rollouts.
- ▶ RecordEpisodeStatistics: This wrapper will keep track of **cumulative rewards** and **episode lengths**.
- ▶ Check at:
  - ▶ [https://gymnasium.farama.org/api/wrappers/misc\\_wrappers/](https://gymnasium.farama.org/api/wrappers/misc_wrappers/)

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake


Bibliography

# Classic RL Examples

## CartPole

The **CartPole** environment:

- ▶ A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track.
- ▶ The pendulum is placed upright on the cart and the goal is to balance the pole by applying forces in the left and right direction on the cart.



# Classic RL Examples

## CartPole

### Observation Space:

- ▶ The **observation** is a `ndarray` with shape `(4,)` with the values corresponding to the following positions and velocities:

Num	Observation	Min	Max
0	Cart Position	-2.4	2.4
1	Cart Velocity	$-\infty$	$+\infty$
2	Pole Angle	$\approx -24^\circ$	$\approx +24^\circ$
3	Pole Angular Velocity	$-\infty$	$+\infty$

# Classic RL Examples

## CartPole

### Action Space:

- ▶ The action is a ndarray with shape (1,) which can take values {0, 1} indicating the direction of the fixed force the cart is pushed with:

Values	Action
0	Push cart to the left
1	Push cart to the right

- ▶ The velocity that is reduced or increased by the applied force is not fixed, and it depends on the angle the pole is pointing.
- ▶ The centre of gravity of the pole varies the amount of energy needed to move the cart underneath it.

# Classic RL Examples

## CartPole

### Rewards:

- ▶ Since the goal is to keep the pole upright for as long as possible, a **reward of +1 for every step taken**, including the termination step, is allotted.
- ▶ The **threshold for rewards** is 500 for v1 and 200 for v0.

# Classic RL Examples

## CartPole

### **Starting State:**

- ▶ All observations are assigned a uniformly random value in (-0.05, 0.05).

# Classic RL Examples

## CartPole

### Episode End:

The episode ends if **any one of the following occurs**:

1. Pole Angle is greater than  $\pm 12^\circ$
2. Cart Position is greater than  $\pm 2.4$  (center of the cart reaches the edge of the display)
3. Episode length is greater than 500 (200 for v0)

# Classic RL Examples

## CartPole

Relevant **questions** regarding this environment:

1. How is the **observation** space?
2. How is the **action** space?
3. What about the **rewards**?

# Classic RL Examples

## CartPole

Relevant **questions** regarding this environment:

1. How is the **observation** space?
  1. **Continuos**, ndarray (4,)
    - ▶ position  $\in (-2.4, 2.4)$
    - ▶ velocity  $\in (-\infty, \infty)$
    - ▶ angle  $\in (-24, 24)$
    - ▶ angular velocity  $\in (-\infty, \infty)$
  2. **Discrete**, i.e.  $\{0, 1\}$
  3. **Discrete**, i.e. +1 for each step “alive”
2. How is the **action** space?
3. What about the **rewards**?

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Classic RL Examples

## FrozenLake

The **FrozenLake** environment:

- ▶ Frozen lake involves **crossing a frozen lake** from **start** to **goal** without falling into any holes by walking over the frozen lake.
- ▶ The player may not always move in the intended direction due to the **slippery nature** of the frozen lake.

S	F	F	F
F	H	F	H
F	F	F	H
H	F	F	G

# Classic RL Examples

## FrozenLake

### Observation Space:

- ▶ The **observation** is an integer value representing the player's current position as:

`current_position = current_row × n_rows + current_col`

where both the row and col start at 0.

Num	Observation	Min	Max
0	Current position	0	$n^2 - 1$

where  $n$  is the size of the grid ( $n \times n$ ).

# Classic RL Examples

## FrozenLake

### Action Space:

- The action shape is  $(1,)$  in the range  $\{0, 3\}$  indicating which direction to move the player.

Values	Action
0	Move left
1	Move down
2	Move right
3	Move up

# Classic RL Examples

## FrozenLake

### Rewards:

- ▶ Reward schedule:

State	Reward
Reach goal	+1
Reach hole	0
Reach frozen	0

# Classic RL Examples

## FrozenLake

### **Starting State:**

- ▶ The episode starts with the **player in state [0]**, i.e. location [0, 0].

# Classic RL Examples

## FrozenLake

### Episode End:

The episode ends if **any one of the following occurs**:

1. **Termination**:

- ▶ The player moves into a hole.
- ▶ The player reaches the goal.

2. **Truncation** (when using the `time_limit` wrapper):

- ▶ The length of the episode is 100 for  $4 \times 4$  environment,
- ▶ or 200 for  $8 \times 8$  environment.

# Classic RL Examples

## FrozenLake

### Information:

1. `step()` and `reset()` return a dict with the following keys:
  - ▶ Transition probability for the state ( $p$ ).
2. `is_slippery` attribute:
  - ▶ If `true` the player will move in intended direction with probability of  $\frac{1}{3}$  else will move in either perpendicular direction with equal probability of  $\frac{1}{3}$  in both directions.

# Classic RL Examples

## FrozenLake

Relevant **questions** regarding this environment:

1. How is the **observation** space?
2. How is the **action** space?
3. What about the **rewards**?

# Classic RL Examples

## FrozenLake

Relevant **questions** regarding this environment:

1. How is the **observation** space?
  2. How is the **action** space?
  3. What about the **rewards**?
1. **Discrete**, integer value  $\in \{0, n^2 - 1\}$ 
    - ▶ where  $n$  is the size of the grid ( $n \times n$ ).
  2. **Discrete**, i.e.  $\{0, 1, 2, 3\}$
  3. **Discrete**, i.e. +1 for reaching the goal state

# Table of Contents

Introduction

Basic structure of an RL scenario

Description

Basic operating diagram

The Gymnasium API

Space

Environment

Wrappers

Classic RL Examples

CartPole

FrozenLake

Bibliography

# Bibliography

## References

Some relevant references:

1. **Farama Foundation**, *Gymnasium Documentation*. Accessed on September 4, 2025.  
<https://gymnasium.farama.org/index.html>